

# COMPSCI 371 Homework 7

**Group Members:** Gordon Liang, Phillip Sievers, Jose Pablo Rivera

## Problem 0 (3 points)

### Part 1: The Gini Index

#### Problem 1.1 (Exam Style)

$$p^* = \begin{pmatrix} 1 \\ 0 \\ .. \\ 0 \end{pmatrix} \quad (1)$$

$$d_{max}^2 = |p^* - c|^2 \quad (2)$$

$$= \left\| \begin{pmatrix} 1 \\ 0 \\ .. \\ 0 \end{pmatrix} - \frac{1}{K} \begin{pmatrix} 1 \\ 1 \\ .. \\ 1 \end{pmatrix} \right\|^2 \quad (3)$$

$$= \left( 1 - \frac{1}{K} \right)^2 + (K-1) \left( -\frac{1}{K} \right)^2 \quad (4)$$

$$= 1 - \frac{2}{K} + \frac{1}{K^2} + \frac{K}{K^2} - \frac{1}{K^2} \quad (5)$$

$$= 1 - \frac{1}{K} \quad (6)$$

#### Problem 1.2 (Exam Style)

$$p = \begin{pmatrix} p_1 \\ p_2 \\ \dots \\ p_K \end{pmatrix} \quad (7)$$

$$d^2(p, c) = |p - c|^2 \quad (8)$$

$$= \left\| \begin{pmatrix} p_1 \\ p_2 \\ \dots \\ p_K \end{pmatrix} - \frac{1}{K} \begin{pmatrix} 1 \\ 1 \\ \dots \\ 1 \end{pmatrix} \right\|^2 \quad (9)$$

$$= \sum_{k=1}^K (p_k - \frac{1}{K})^2 \quad (10)$$

$$= \sum_{k=1}^K p_k^2 - \sum_{k=1}^K 2 \frac{p_k}{K} + \sum_{k=1}^K \frac{1}{K^2} \quad (11)$$

$$= \sum_{k=1}^K p_k^2 - 2 \frac{1}{K} + \frac{1}{K} \quad (12)$$

$$= \sum_{k=1}^K p_k^2 - \frac{1}{K} \quad (13)$$

### Problem 1.3 (Exam Style)

$$d_{max}^2 - d^2(p, c) = 1 - \frac{1}{K} - \left( \sum_{k=1}^K p_k^2 - \frac{1}{K} \right) \quad (14)$$

$$= 1 - \sum_{k=1}^K p_k^2 \quad (15)$$

$$= i_{Gini}(p) \quad (16)$$

## Part 2: Basics of Decision Tree Classifiers

### Problem 2.1 (Exam Style)

e

1

0.3

0

0.320

0.2

### Problem 2.2 (Exam Style)

In [550...]

```
import numpy as np
import matplotlib.pyplot as plt
```

In [551...]

```
class Node:
    def __init__(self, j, t):
        self.j = j
        self.t = t
        self.l = None
        self.r = None
        self.lp = None
        self.rp = None

    a = Node(1,6)
    b = Node(2,3)
    c = Node(2,1)
    d = Node(1,4)
    e = Node(2,7)

    a.l = b
    b.l = c
    b.r = d
    a.r = e

    c.lp = (0.3,0.4,0.3)
    c.rp = (0.2,0.2,0.6)

    d.lp = (0.0,0.2,0.8)
    d.rp = (0.6,0.1,0.3)

    e.lp = (0.1,0.5,0.4)
    e.rp = (0.8,0.1,0.1)
```

In [552...]

```
def h(n, x):
    if not n.l and not n.r:
        if x[n.j - 1] <= n.t:
            return np.argmax(n.lp) + 1
        return np.argmax(n.rp) + 1
    if x[n.j - 1] <= n.t:
        return h(n.l, x)
    else:
        return h(n.r, x)


xx = np.arange(0,11)
yy = np.arange(0,11)
arr = [[h(a, (x,y)) for x in xx] for y in yy]

img = plt.imshow(
    arr,
    cmap='terrain',
    origin='lower',
    extent=[-1, 10, -1, 10],
    interpolation='nearest'
)

plt.text(2,9,3)
plt.text(8,9,1)
plt.text(8,4,2)

plt.xlim(0, 10)
```

```
plt.ylim(0, 10)  
plt.show()
```



### Problem 2.3 (Exam Style)

$$i(S) = 3/5$$

$j$	$t$	$\ L\ $	$i_L$	$\ R\ $	$i_R$	$\delta$	<b>best</b>
1	3	2	0	3	1/3	2/5	yes
1	5	3	1/3	2	1/2	1/5	no
1	7	4	1/2	1	0	1/5	no
2	2	1	0	4	1/2	1/5	yes
2	4	2	1/2	3	1/3	1/5	yes
2	6	3	2/3	2	1/2	0	no

## Part 3: Regression Trees

In [553]:

```
import pickle  
import numpy as np  
from types import SimpleNamespace  
from matplotlib import pyplot as plt  
from matplotlib.ticker import MultipleLocator  
from matplotlib.patches import Rectangle  
from matplotlib import use  
from matplotlib import colors, colormaps  
plt.style.use('default')  
%matplotlib inline
```

```
In [554...]def bounding_box(xs, margin=0.05):
    mn, mx = np.min(xs, axis=0) - margin, np.max(xs, axis=0) + margin
    return SimpleNamespace(left=mn[0], right=mx[0], bottom=mn[1], top=mx[1])
```

```
In [555...]def replace_side(box, side, value):
    new = dict(**box.__dict__)
    new[side] = value
    return SimpleNamespace(**new)
```

```
In [556...]def shade_box(box, color):
    corner = box.left, box.bottom
    width, height = box.right - corner[0], box.top - corner[1]
    rectangle = Rectangle(corner, width, height,
                           edgecolor='none', facecolor=color)
    plt.gca().add_patch(rectangle)
```

```
In [557...]def make_colormap(data):
    value_range = [np.minimum(0, np.min(data.y)), np.maximum(0, np.max(data.y))]
    half_range = max(-value_range[0], value_range[1])
    return SimpleNamespace(
        map=colormaps['RdYlBu'],
        norm=colors.CenteredNorm(vcenter=0, halfrange=half_range)
    )
```

```
In [558...]def value_color(value, color_scheme):
    return np.squeeze(color_scheme.map(color_scheme.norm(value)))
```

```
In [559...]def make_figure(box=None, fig_size=(6, 6)):
    if box is None:
        box = SimpleNamespace(left=0, right=1, bottom=0, top=1)
    plt.figure(figsize=fig_size, tight_layout=True)
    plt.plot((box.left, box.right, box.right, box.left, box.left),
              (box.bottom, box.bottom, box.top, box.top, box.bottom),
              lw=0.5, c='gray')
    plt.axis('square')
    plt.axis('off')
    return box
```

```
In [560...]import urllib.request
import ssl
from os import path as osp
import shutil

# same as zsh
# curl https://www2.cs.duke.edu/courses/fall25/comsci371/homework/7/regressor_data.pkl
# --output regressor_data.pkl
def retrieve(file_name, semester='fall25', homework=7):
    if osp.exists(file_name):
        print('Using previously downloaded file {}'.format(file_name))
    else:
        context = ssl._create_unverified_context()
        fmt = 'https://www2.cs.duke.edu/courses/{}/comsci371/homework/{}/{}'
        url = fmt.format(semester, homework, file_name)
        with urllib.request.urlopen(url, context=context) as response:
            with open(file_name, 'wb') as file:
                shutil.copyfileobj(response, file)
        print('Downloaded file {}'.format(file_name))
```

```
In [561... regression_training_set_name = 'regressor_data.pkl'  
retrieve(regression_training_set_name)
```

Using previously downloaded file regressor\_data.pkl

```
In [562... with open(regression_training_set_name, 'rb') as file:  
    regression_training_set = pickle.load(file)
```

```
In [563... color_map = make_colormap(regression_training_set)
```

```
In [564... regression_training_set  
X = regression_training_set.x  
y = regression_training_set.y
```

```
In [565... color_map = make_colormap(regression_training_set)
```


```
In [566... def plot_data(data, color_scheme):  
    plt.scatter(  
        data.x[:, 0], data.x[:, 1], s=10,  
        c=data.y, cmap=color_scheme.map, norm=color_scheme.norm  
    )  
    plt.colorbar(cmap=color_scheme.map, norm=color_scheme.norm, shrink=0.855)
```

```
In [567... import sklearn  
h = sklearn.tree.DecisionTreeRegressor()  
h.fit(X,y)
```

Out[567... ▾ DecisionTreeRegressor ⓘ ⓘ

► Parameters

```
In [568... def draw_tree(t, b, color_scheme):  
    def draw(node, box):  
        if t.children_left[node] == -1 \  
        and t.children_right[node] == -1:  
            shade_box(box, value_color(t.value[node], color_scheme))  
            return  
        plt.plot([(t.threshold[node], t.threshold[node]),(box.left, box.right)][t.feature  
                           [(box.bottom, box.top),(t.threshold[node], t.threshold[node])][t.feature  
                               lw = 0.25, c='gray')  
        draw(t.children_left[node], replace_side(box,['right','top'][t.feature[node]],t.  
        draw(t.children_right[node], replace_side(box,['left','bottom'][t.feature[node]]  
  
        draw(0, b)  
        plot_data(regression_training_set, color_map)  
  
    data_box = make_figure()  
    draw_tree(h.tree_, data_box, color_map)
```



## Part 4: Random Decision Forests


```
In [569...]: spiral_file = 'spiral.pkl'  
retrieve(spiral_file)  
with open(spiral_file, 'rb') as file:  
    spiral = pickle.load(file)
```

Using previously downloaded file spiral.pkl

```
In [570...]: tx = spiral.train.x  
ty = spiral.train.y  
  
sx = spiral.test.x  
sy = spiral.test.y
```

```
In [571...]: def draw_spiral(s):  
    font_size = 14  
    plt.figure(figsize=(5, 5), tight_layout=True)  
    for m, color in enumerate(s.labels):  
        select = s.train.y == m  
        plt.scatter(s.train.x[select, 0],  
                    s.train.x[select, 1],  
                    c=color)  
    plt.gca().set_aspect(1)  
    plt.axis('off')  
    plt.title('spiral data', fontsize=font_size)  
    plt.draw()  
  
draw_spiral(spiral)
```

## spiral data



```
In [572]: from matplotlib.colors import ListedColormap
```

```
def coarse_regions(h, colors, step=0.01):
    xx, yy = np.meshgrid(np.arange(0, 1, step),
                         np.arange(0, 1, step))
    label = h.predict(np.c_[xx.ravel(), yy.ravel()]).reshape(xx.shape)
    color_map = ListedColormap(colors)
    plt.contourf(xx, yy, label, cmap=color_map)
```

```
In [573]: h = sklearn.tree.DecisionTreeClassifier()
h.fit(tx,ty)
```

```
Out[573]:
```

```
► Parameters
```

### Problem 4.1

```
In [574]: import warnings
```

```
def addplot(n_estimators, subplot):
    h = sklearn.ensemble.RandomForestClassifier(
        n_estimators = n_estimators,
        max_depth = None,
        min_samples_split = 2,
        random_state = 0,
        oob_score = True
    )
```

```

h.fit(tx,ty)

plt.subplot(1,2,subplot)
coarse_regions(h, spiral.labels)

print(f'accuracy for {n_estimators} trees')
print(np.round(sklearn.metrics.accuracy_score(sy,h.predict(sx)) * 100,3),'%')
print(f'out of bag score for {n_estimators} trees')
print(np.round(h.oob_score_* 100,3),'%')

warnings.catch_warnings()
warnings.simplefilter('ignore')

addplot(5,1)
addplot(500,2)
plt.show()

```

accuracy for 5 trees

91.125 %

out of bag score for 5 trees



65.25 %

accuracy for 500 trees

96.45 %

out of bag score for 500 trees

90.25 %



## Problem 4.2 (Exam Style)

1. Why is the OOB accuracy off when 5 trees are used? Be quantitative, and refer to the slides (not the notes) on trees and random forests for this part.

The number of trees used to validate each OOB data point follows a binomial distribution

$$X \sim B(p = 0.37, N = 5)$$

As demonstrated in the calculations and distribution graph below, this means that the number of trees used to validate any given OOB data point is quite low.

For instance,  $p(X = 1) = 0.291$  and  $p(X = 2) = 0.342$ .

This reduces the number of "votes" per data point and biases the prediction to a few trees'  $B_m$ .

Furthermore,  $p(X = 0) = 0.099$  just means there is a 9.9% chance that  $x$  is not in  $T'$  at all.

This reduces the size of  $T'$  and may further reduce how meaningful the OOB is.

2. How good is the OOB accuracy estimate with 500 trees? Compare it with the test accuracy.

OOB = 90.25%

test accuracy = 96.45%

OOB accuracy is a good but not great estimate. It is meaningfully (6.20%) under test accuracy.

3. How do the two sets of decision regions (obtained with 5 and 500 trees) and the corresponding test accuracy values compare to each other, qualitatively?

- approximately equal spiral appearance of decision boundaries
- $M = 5$  has rougher decision boundaries
- $M = 500$  has smoother decision boundaries
- accuracy for  $M = 5$  is 91.125%
- accuracy for  $M = 500$  is 96.45%
- more trees is more accurate
- thus, smoother decision boundaries represent data better

In [575...]

```
# 4.2.1



import math

cc = np.arange(6)
freq = [math.comb(5, c) * 0.37**c * 0.63**(5-c) for c in cc]
cdf = np.cumsum(freq)

for c in range(6):
    print(f'Probability {c} trees vote on any given sample', np.round(freq[c],3))

fig, ax = plt.subplots(2,1)
fig.subplots_adjust(hspace=0.5)
ax[0].bar(cc, freq)
ax[0].set_title('Probability k out of 5 trees vote on any given sample')
ax[1].bar(cc, cdf)
ax[1].set_title('Cumulative Probability k out of 5 trees vote on any given sample')
plt.show()
```

```
Probability 0 trees vote on any given sample 0.099
Probability 1 trees vote on any given sample 0.291
Probability 2 trees vote on any given sample 0.342
Probability 3 trees vote on any given sample 0.201
Probability 4 trees vote on any given sample 0.059
Probability 5 trees vote on any given sample 0.007
```



## Part 5: Trees vs. Forests

```
In [576...]: mnist_file_name = 'full_mnist.pkl'
retrieve(mnist_file_name)
```

Using previously downloaded file full\_mnist.pkl

```
In [577...]: with open(mnist_file_name, 'rb') as file:
    digits = pickle.load(file)
tx = digits.train.x
ty = digits.train.y
sx = digits.test.x
sy = digits.test.y
```


```
In [578...]: def x_to_image(x, stats):
    x = np.round(x * stats.std + stats.mean)
    x = np.clip(x * 255. / stats.max, 0., 255.).astype(np.uint8)
    return np.reshape(x, stats.shape)

def show_random_training_images(data, rows=3, columns=6):
    xs, ys = data.train.x, data.train.y
    rng = np.random.default_rng()
    indices = rng.integers(low=0, high=len(ys), size=rows * columns)
    plt.figure(figsize=(1.8 * columns, 1.9 * rows), tight_layout=True)
    for plot, index in enumerate(indices):
        image = x_to_image(xs[index], data.stats)
        plt.subplot(rows, columns, plot + 1)
        plt.imshow(image, cmap='gray')
```

```

plt.axis('off')
plt.title(ys[index], fontsize=18)
plt.show()
show_random_training_images(digits)

```



## Problem 5.1

```

In [579... def display_stats(h):
    print(f'train error rate: {np.round((1-h.score(tx,ty))*100,3)}%')
    print(f'test error rate: {np.round((1-h.score(sx,sy))*100,3)}%')

    confusion_matrix = sklearn.metrics.confusion_matrix(sy,h.predict(sx))
    print('Confusion Matrix')
    print(confusion_matrix)

    fig, ax = plt.subplots()
    ax.xaxis.set_major_locator(MultipleLocator(1))
    ax.yaxis.set_major_locator(MultipleLocator(1))
    ax.set_title('Confusion Matrix on MNIST data')
    ax.imshow(confusion_matrix, cmap='turbo')
    plt.show()

```

```

In [580... h = sklearn.tree.DecisionTreeClassifier()
h.fit(tx,ty)

```

Out [580... ▾ DecisionTreeClassifier ⓘ ⓘ

► Parameters

```

In [581... display_stats(h)


```

train error rate: 0.0%  
test error rate: 12.43%

#### Confusion Matrix

[ [ 916	0	9	7	5	11	10	4	10	8]
[ 0	1089	8	6	2	6	7	3	13	1]
[ 10	11	877	32	14	14	13	25	26	10]
[ 8	7	29	858	4	54	4	3	21	22]
[ 5	3	11	7	858	11	15	9	21	42]
[ 12	9	5	44	7	741	25	4	24	21]
[ 17	3	10	8	17	16	848	4	25	10]
[ 5	15	23	26	6	4	1	914	7	27]
[ 7	6	23	36	26	30	20	10	791	25]
[ 15	2	7	17	42	9	6	22	24	865]]

Confusion Matrix on MNIST data



## Problem 5.2 (Exam Style)

1. Give an example of a pair of different digits that your confusion matrix indicates are often confused with each other. You need not display the digits.

5 is often mistaken as a 3

4 is often mistaken as a 9

2. Is the confusion always symmetric in the matrix? Symmetry here means that for any digit pair it is about equally likely to mistake the first digit for the second as it is to mistake the second for the first. No need to justify your answer.

no

3. Does the decision tree underfit? How can you tell, and what is the reason?

no. achieves 0.0% error on training data.

DecisionTreeClassifier with default parameters is a good model for this task, especially if it has seen the data.

4. Does the decision tree overfit? How can you tell, and what might be the reason?

yes. 12.33% error on testing data is much greater than 0.0% error on training data.

DecisionTreeClassifier with default parameters is not as good for classifying new data.


## Problem 5.3

```
In [582...]  
h = sklearn.ensemble.RandomForestClassifier(  
    n_estimators = 100,  
    oob_score = True  
)  
  
h.fit(tx,ty)
```

Out [582...]  
▼ RandomForestClassifier ⓘ ⓘ  
► Parameters

```
In [583...]  
display_stats(h)  
print(f'out of bag error rate: {np.round((1-h.oob_score_)*100,3)}%')  
  
train error rate: 0.0%  
test error rate: 2.99%  
Confusion Matrix  
[[ 971   0   0   0   0   1   2   2   4   0]  
 [   0 1123   3   3   0   2   2   0   1   1]  
 [   6   0 1002   6   1   0   4   9   4   0]  
 [   0   0   8  972   0  10   0   9   8   3]  
 [   1   0   1   0  956   0   6   0   2  16]  
 [   3   0   1  10   3  857   7   2   5   4]  
 [   5   3   0   0   5   4  937   0   4   0]  
 [   1   5  20   2   1   0   0  987   1  11]  
 [   4   0   4   8   3   6   4   3  934   8]  
 [   5   6   0  11  11   4   1   3   6  962]]
```

Confusion Matrix on MNIST data



out of bag error rate: 3.452%

### Problem 5.4 (Exam Style)

1. Does the random forest underfit?

no

train error rate: 0.0%

test error rate: 2.99%

2. Does the random forest overfit more or less than the decision tree you found earlier for this dataset?

random forest overfits less than the decision tree

RF error rate 2.95% < DT error rate 12.33%

3. Is the out-of-bag error rate a reasonably accurate estimate (say, within 1 or 2 percentage points) of the test error rate?

yes, OOB error rate is 3.468%, which is within a few percentage points from the test error rate 2.95%